


Image Greyscale Benchmark Results

Input Image

- Image dimensions (pixels): 28000x28000 (FHD == 1920x1080)



- Image size (bytes): 119,828,615

Hardware/Software Specifications

Hardware

CPU

CPU: Info: Quad Core
model: Intel Core i7-4790K
bits: 64
type: MT MCP
arch: Haswell
rev: 3
cache:
L2: 8 MiB
flags: avx avx2 lm nx pae sse sse2 sse3 sse4_1 sse4_2 ssse3 vmx
bogomips: 63873
Speed: 4333 MHz
min/max: 800/4400 MHz
Core speeds (MHz):
1: 4333 2: 4043 3: 4323 4: 4126
5: 4346 6: 4282 7: 4277 8: 4339

RAM

Memory: RAM: total: 11.54 GiB
Array-1: capacity: 32 GiB
slots: 4
EC: None
max-module-size: 8 GiB
note: est.
Device-1: ChannelA-DIMM0 size: No Module Installed
Device-2: ChannelA-DIMM1 size: 4 GiB speed: 1333 MT/s type: DDR3
Device-3: ChannelB-DIMM0 size: No Module Installed
Device-4: ChannelB-DIMM1 size: 8 GiB speed: 1333 MT/s type: DDR3

CUDA

CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0:
CUDA Driver Version / Runtime Version 11.3 / 11.3
CUDA Capability Major/Minor version number: 5.0
Total amount of global memory: 3998 MBytes (4192337920 bytes)
(005) Multiprocessors, (128) CUDA Cores/MP: 640 CUDA Cores
GPU Max Clock rate: 1110 MHz (1.11 GHz)
Memory Clock rate: 2700 Mhz
Memory Bus Width: 128-bit
L2 Cache Size: 2097152 bytes
Total amount of constant memory: 65536 bytes
Total amount of shared memory per block: 49152 bytes
Total shared memory per multiprocessor: 65536 bytes
Total number of registers available per block: 65536
Warp size: 32
Maximum number of threads per multiprocessor: 2048
Maximum number of threads per block: 1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)

deviceQuery
CUDA Driver = CUDART
CUDA Driver Version = 11.3
CUDA Runtime Version = 11.3
NumDevs = 1
Result = PASS

Software

Operating System

Kernel: 5.12.14 x86_64

Compilers and Tools

- gcc:
gcc (GCC) 11.1.0
- clang:
clang version 12.0.1
Target: x86_64-pc-linux-gnu
Thread model: posix
- cmake:
cmake version 3.20.5

Algorithm

The overall algorithm of all implementation models is as the following:

- Read the input image
- Load the entire image into memory (RAM or GPU DRAM)
- Start timer
- Transform image (launch threads in multi-threaded implementation)
- Gather results
- Stop timer
- Print elapsed execution time
- Write results to output image file
- Clean any allocated memory


Limitations

Current implementation uses libjpeg for reading, manipulating, and writing jpeg images.


Results

Results are average of *only processing time* of executing the same implementation with different configuration parameters 2000 times:


- OpenMP



- pthread



- CUDA



Speedup

The following is the best execution time of different parallelism model implementations compared in one view (+ single-threaded serial execution):

